

# Group A

## Assignment 2

### Data Wrangling II

Import all the required Python Libraries.

```
# code here
import pandas as pd
import numpy as np
import random
import seaborn as sns
```

Create a DataFrame from the dictionary

```
# code here
data={"Stud_id":[i for i in range(1,101)],
      "Gender":[random.choice(['Female','Male']) for _ in
range(100) ],
      "Class":[random.choice(['SE','TE','BE']) for _ in range(100) ],
      "Attendance":[random.uniform(0,100) for _ in range(100)],
      "DSBDA":[random.randint(0,100) for _ in range(100)],
      "WT":[random.randint(0,100) for _ in range(100)],
      "AI":[random.randint(0,100) for _ in range(100)]}
```

data

```
df=pd.DataFrame(data)
```

df

	Stud_id	Gender	Class	Attendance	DSBDA	WT	AI
0	1	Male	BE	30.239580	56	82	7
1	2	Female	SE	44.727962	63	49	34
2	3	Female	SE	53.430196	60	65	45
3	4	Male	BE	32.110630	91	69	98
4	5	Female	BE	84.151293	94	46	66
...	...	...	...	...	...	...	...
95	96	Female	SE	82.940270	35	28	93
96	97	Male	TE	27.447840	88	82	49
97	98	Male	TE	93.815668	15	74	53
98	99	Female	TE	8.431059	94	81	71
99	100	Female	SE	62.632263	77	71	20

[100 rows x 7 columns]

```
df.to_csv('AcademicPerformance.csv')
```

Load the Dataset into pandas dataframe.

# [code here](#)

```
Std_df=pd.read_csv('StudentPerformance.csv')
```

```
Std_df.head()
```

	gender	race/ethnicity	parental level of education	lunch
0	female	group B	bachelor's degree	standard
1	female	group C	some college	standard
2	female	group B	master's degree	standard
3	male	group A	associate's degree	free/reduced
4	male	group C	some college	standard

	test preparation course	math score	reading score	writing score
0	none	72.0	72.0	74.0
1	completed	69.0	90.0	88.0
2	none	90.0	95.0	93.0
3	none	47.0	57.0	44.0
4	none	76.0	78.0	75.0

```
Std_df.shape
```

```
(1000, 8)
```

```
Std_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	gender	985 non-null	object
1	race/ethnicity	989 non-null	object
2	parental level of education	993 non-null	object
3	lunch	999 non-null	object
4	test preparation course	990 non-null	object
5	math score	990 non-null	float64
6	reading score	990 non-null	float64
7	writing score	990 non-null	float64

```
dtypes: float64(3), object(5)
```

```
memory usage: 62.6+ KB
```

```
Std_df.describe()
```

	math score	reading score	writing score
count	990.000000	990.000000	990.000000
mean	66.055556	69.116162	68.082828
std	15.137922	14.594195	15.158456
min	0.000000	17.000000	10.000000
25%	57.000000	59.000000	58.000000
50%	66.000000	70.000000	69.000000

75%	77.000000	79.000000	79.000000
max	100.000000	100.000000	100.000000

## Data Preprocessing

Scan all variables for missing values and inconsistencies. If there are missing values and/or inconsistencies, use any of the suitable techniques to deal with them.

```
# code here
Std_df.isnull().sum()

gender          15
race/ethnicity  11
parental level of education  7
lunch           1
test preparation course  10
math score      10
reading score   10
writing score   10
dtype: int64

Std_df.nunique()

gender          2
race/ethnicity  5
parental level of education  6
lunch           2
test preparation course  2
math score      81
reading score   72
writing score   77
dtype: int64

Std_df['gender'].value_counts()

gender
female    510
male      475
Name: count, dtype: int64

Std_df.gender.fillna('female',inplace=True)

Std_df.isnull().sum()

gender          0
race/ethnicity  11
parental level of education  7
lunch           1
test preparation course  10
math score      10
reading score   10
```

```
writing score          10
dtype: int64

Std_df['race/ethnicity'].value_counts()

race/ethnicity
group C    315
group D    259
group B    189
group E    139
group A     87
Name: count, dtype: int64

Std_df['race/ethnicity'].fillna('group C',inplace=True)

Std_df.isnull().sum()

gender          0
race/ethnicity  0
parental level of education  7
lunch           1
test preparation course  10
math score      10
reading score   10
writing score    10
dtype: int64

Std_df['parental level of education'].value_counts()

parental level of education
some college    224
associate's degree  221
high school    196
some high school  178
bachelor's degree  116
master's degree   58
Name: count, dtype: int64

Std_df['parental level of education'].fillna('some college',inplace=True)

Std_df.isnull().sum()

gender          0
race/ethnicity  0
parental level of education  0
lunch           1
test preparation course  10
math score      10
reading score    10
```

```
writing score          10
dtype: int64

Std_df['lunch'].value_counts()

lunch
standard      644
free/reduced   355
Name: count, dtype: int64

Std_df.lunch.fillna('standard',inplace=True)

Std_df.isnull().sum()

gender          0
race/ethnicity  0
parental level of education  0
lunch           0
test preparation course  10
math score      10
reading score   10
writing score   10
dtype: int64

Std_df['test preparation course'].value_counts()

test preparation course
none      632
completed 358
Name: count, dtype: int64

Std_df['test preparation course'].fillna('none',inplace=True)


Std_df.isnull().sum()

gender          0
race/ethnicity  0
parental level of education  0
lunch           0
test preparation course  0
math score      10
reading score   10
writing score   10
dtype: int64
```

Using Seaborn lib to plot graph

```
sns.kdeplot(Std_df['math score'])

<Axes: xlabel='math score', ylabel='Density'>
```



```
sns.distplot(Std_df['reading score'])
```

C:\Users\Ashitosh Warghade\AppData\Local\Temp\ipykernel\_14772\60741464.py:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(Std_df['reading score'])
```

<Axes: xlabel='reading score', ylabel='Density'>



```
sns.distplot(Std_df['writing score'],hist=False)
```

C:\Users\Ashitosh Warghade\AppData\Local\Temp\ipykernel\_14772\3580858108.py:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(Std_df['writing score'],hist=False)
```

<Axes: xlabel='writing score', ylabel='Density'>



```
sns.distplot(Std_df['writing score'],kde=False)
```

C:\Users\Ashitosh Warghade\AppData\Local\Temp\ipykernel\_14772\2837882268.py:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(Std_df['writing score'],kde=False)
```

<Axes: xlabel='writing score'>





```
Std_df["math score"].skew()
```

```
-0.27756569959735494
```

```
Std_df["math score"].fillna(Std_df["math score"].mean(), inplace=True)
```

```
Std_df.isnull().sum()
```

gender	0
race/ethnicity	0
parental level of education	0
lunch	0
test preparation course	0
math score	0
reading score	10
writing score	10
dtype: int64	

```
Std_df["writing score"].mode()
```

```
0    74.0
```

```
Name: writing score, dtype: float64
```

```
Std_df["writing score"].mode()[0]
```

```
74.0
```

```
Std_df["writing score"].fillna(Std_df["writing score"].mode()[0],  
inplace=True)
```

```
Std_df.isnull().sum()
```

gender	0
race/ethnicity	0
parental level of education	0
lunch	0
test preparation course	0
math score	0
reading score	10
writing score	0
dtype: int64	

```
Std_df["reading score"].fillna(Std_df["reading score"].median(),  
inplace=True)
```


```
Std_df.isnull().sum()
```

gender	0
race/ethnicity	0
parental level of education	0
lunch	0
test preparation course	0
math score	0
reading score	0
writing score	0
dtype: int64	


Scan all numeric variables for outliers. If there are outliers, use any of the suitable techniques to deal with them.

```
# code here  
sns.boxplot(Std_df)
```

```
<Axes: >
```



```
def treat_outlier(df,col_list):  
    for col_name in col_list:  
        q1=df[col_name].quantile(0.25)  
        q3=df[col_name].quantile(0.75)  
        iqr=q3-q1  
        lower=q1-(1.5*iqr)  
        upper=q3+(1.5*iqr)  
        np.clip(df[col_name],lower,upper,inplace=True)  
  
treat_outlier(Std_df,['math score','writing score','reading score'])  
sns.boxplot(Std_df)  
<Axes: >
```



Apply data transformations on at least one of the variables. The purpose of this transformation should be one of the following reasons: to change the scale for better understanding of the variable, to convert a non-linear relation into a linear one, or to decrease the skewness and convert the distribution into a normal distribution.

```
# code here
Std_df["math score"].skew()
-0.12912399951580147

Std_df["writing score"].skew()
-0.24169581899585696

Std_df["reading score"].skew()
-0.20943942810853586

Std_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   gender              1000 non-null   object
```

1	race/ethnicity	1000	non-null	object
2	parental level of education	1000	non-null	object
3	lunch	1000	non-null	object
4	test preparation course	1000	non-null	object
5	math score	1000	non-null	float64
6	reading score	1000	non-null	float64
7	writing score	1000	non-null	float64

dtypes: float64(3), object(5)

memory usage: 62.6+ KB

*# when we don't know the data*

*# data normally distribute then use standard scale(mostly -3 to 3)*

*# data is not normally then min max scale(0-1)*

```
from sklearn.preprocessing import MinMaxScaler , StandardScaler
```

*# object scaler, class StandardScaler*

```
scaler=StandardScaler()
```

*# to scaler to learn the data frame math score col pattern*

```
scaler.fit(Std_df[['math score']])
```

```
StandardScaler()
```

```
scaled_mathscore= scaler.transform(Std_df[['math score']])
```

```
scaled_mathscore
```

*# to insert we need 3 parameter 1st col index,col name,inserting val*

```
Std_df.insert(6,'scaled math score',scaled_mathscore)
```

```
Std_df.head()
```

	gender	race/ethnicity	parental level of education	lunch	\
0	female	group B	bachelor's degree	standard	
1	female	group C	some college	standard	
2	female	group B	master's degree	standard	
3	male	group A	associate's degree	free/reduced	
4	male	group C	some college	standard	

	test preparation course	math score	scaled math score	reading score	\
0	none	72.0	0.396213	72.0	
1	completed	69.0	0.193129	90.0	
2	none	90.0	1.614718	95.0	
3	none	47.0	-1.296154	57.0	

4	none	76.0	0.666992
---	------	------	----------

78.0

	writing score
0	74.0
1	88.0
2	93.0
3	44.0
4	75.0

```
scaler.fit(Std_df[['reading score']])
```

```
StandardScaler()
```

```
scaled_readscore= scaler.transform(Std_df[['reading score']])
```

```
Std_df.insert(8,'scaled reading score',scaled_readscore)
```

```
Std_df.head()
```

	gender	race/ethnicity	parental level of education	lunch	\
0	female	group B	bachelor's degree	standard	
1	female	group C	some college	standard	
2	female	group B	master's degree	standard	
3	male	group A	associate's degree	free/reduced	
4	male	group C	some college	standard	

	test preparation course	math score	scaled math score	reading score \
0	none	72.0	0.396213	72.0
1	completed	69.0	0.193129	90.0
2	none	90.0	1.614718	95.0
3	none	47.0	-1.296154	57.0
4	none	76.0	0.666992	78.0

	scaled reading score	writing score
0	0.197199	74.0
1	1.445736	88.0
2	1.792552	93.0
3	-0.843248	44.0
4	0.613378	75.0

```
scaler=MinMaxScaler()
```

```
scaler.fit(Std_df[['writing score']])
```

```
MinMaxScaler()
```

```
scaled_writescore= scaler.transform(Std_df[['writing score']])
```

```
Std_df.insert(10,'scaled writing score',scaled_writescore)
```

```
Std_df.tail()
```

	gender	race/ethnicity	parental level of education	lunch \
995	female	group E	master's degree	standard
996	male	group C	high school	free/reduced
997	female	group C	high school	free/reduced
998	female	group D	some college	standard
999	female	group D	some college	free/reduced

	test preparation course	math score	scaled math score	reading score \
995	completed	88.0	1.479328	99.0
996	none	62.0	-0.280734	55.0
997	completed	59.0	-0.483818	71.0
998	completed	68.0	0.125434	78.0
999	none	77.0	0.734687	86.0

	scaled reading score	writing score	scaled writing score
995	2.070004	95.0	0.931973
996	-0.981974	55.0	0.387755
997	0.127836	65.0	0.523810
998	0.613378	77.0	0.687075
999	1.168284	86.0	0.809524